

Homework 1
Released: Sep 26, 2024
Due date: Oct 3, 2024, 11:59pm
Total points: 100

Problem 1 (10 points) What is the meaning of the term busy-waiting? What other kinds of waiting are there in an operating system? Can busy waiting be avoided altogether? Explain your answers.

Answer:

Busy waiting:

When thread A requires resources(like lock, I/O) which is kept by another thread B. And thread A keep continuous checking in a loop. Which will cause a busy waiting, because the check instructions consume the CPU resources, CPU can not switch to other instructions. Like Spinlock

Other kind of waiting:

Sleep waiting: When a thread waiting for resources, instead of consuming CPU time, the thread is put to sleep; and the cpu can be allocated to others. like semaphores, mutex with sleep-wait behaviors.

Interrupt/Signal: the process or kernel sets up an interrupt or signal to notify it when the event happened. Which is commonly used in devices drivers, the hardware notifies the CPU that an operation is complete.

Solve busy waiting:

Busy-waiting can be minimized or avoid through more efficient synchronization techniques. Like semaphores, mutex with sleep-wait behaviors.

Problem 2 (10 points) Explain why implementing synchronization primitives by disabling interrupts is not appropriate in a single-processor system if the synchronization primitives are to be used in user-level programs.

Answer:

1, Disabling interrupts in user-level programs may cause the system hangs when an infinite loop or deadlock occurs in the program.

2, Even if the user-level program execute correctly, disabling interrupts will cause the system to be unable to serve other requests, which will slow down the system, especially when the program need access I/O devices.

3, Not safe

Problem 3 (10 points) Given the following program that uses three memory segments in an address space as described in class (code segment, data segment, and stack segment):

```
char a[100];
main(int argc, char ** argv)
{
    int d;
    static double b;
    char *s = "boo", * p;

    p = malloc(300);
    return 0;
}
```

Identify the segment in which each variable resides and indicate if the variable is private to the thread or is shared among threads.

Answer:

variable	Which segment?	Shared among threads?	Reasons
a	Data segment	shared	Global variable
b	Data segment	shared	Static variable
d	Stack segment	private	Local variable within functions
char *s= 'boo'	Pointer s is in Stack segment "boo" is in Code segment	Pointer s is private "boo" is shared	Code segment is shared by threads
char *p=malloc(300)	Pointer p is in Stack segment The memory allocated by malloc() is located in Heap segment	Point p is private The memory allocated in the heap is shared	

Problem 4 (10 points) Which of the following instructions should be privileged:

Answer:

Instructions	Privileged
a. Set the value of a timer	Yes
b. Read the clock	No
c. Clear memory	Yes
d. Issue a trap instruction	No
e. Turn off interrupts	Yes
f. Modify entries in device-status table	Yes
g. Switch from user to kernel mode	Yes
h. access I/O device	Yes

Problem 5 (10 points) Given the following piece of code:

```
main(int argc, char ** argv)
```

```
{
    int child = fork();
    int c = 5;

    if(child == 0)
    {
        c += 5;
    }
    else
    {
        child = fork();
        c += 10;
        if(child)
        {
            c += 5;
        }
    }
}
```

How many different copies of the variable *c* are there? What are their values?

Answer:

There are three different copies of variable *c*, in three different process:

In parent process : *c*=20

In child process 1 create by first `fork()` : *c*=10

In child process 2 create by second `fork()` : *c*=10

Problem 6 (10 points) Assume that a system has multiple processing cores. For each of the following scenarios, describe which is a better locking mechanism - a spinlock or a mutex lock where waiting processes sleep while waiting for the lock to become available:

- The lock is to be held for a short duration.
- The lock is to be held for a long duration.
- A thread may be put to sleep while holding the lock.

Answer:

(1) The lock is to be held for a short duration

The spinlock is a better choice, because spinlock is a busy-waiting lock that can keep the thread running in CPU.

For mutex lock, because put thread to sleep or wake up need context switch, if only wait for a short duration, these overhead may be more costly.

(2) The lock is to be held for a long duration

Mutex lock is better choice, because if use spinlock, keeping thread busy waiting in CPU for a long duration is costly.

(3) A thread may be put to sleep while holding the lock

Mutex lock is better choice, the reason is similar to (2), busy waiting will cause high CPU usage.

Problem 7 (10 points) A multithreaded web server wishes to keep track of the number of requests it services (known as *hits*). Consider the two following strategies to prevent a race condition on the variable hits. The first strategy is to use a basic mutex lock when updating hits:

```
int hits;  
mutex lock hit lock;  
hit lock.acquire();  
hits++;  
hit lock.release();
```

A second strategy is to use an atomic integer:

```
atomic_t hits;  
atomic_inc(&hits);
```

Explain which of these two strategies is more efficient.

Answer:

The second strategy is more efficient.

Because both of them can protect the variable hits, but when use mutex lock, the program need acquire and release lock, which need involves system calls and lead to significant overhead. For atomic integers, it can minimize the overhead associated with locking and unlocking mutexes

Problem 8 (10 points) Servers can be designed to limit the number of open connections. For example, a server may wish to have only N socket connections at any point in time. As soon as N connections are made, the server will not accept another incoming connection until an existing connection is released. Illustrate how semaphores can be used by a server to limit the number of concurrent connections.

Answer:

- 1, When initiate a semaphore, we can set a initial count of N .
 - 2, Before the server accepting a new connection, it will check the semaphore, if the value is greater than 0, which means there is an available slot for a new connection. The server should decrement the semaphore, and the new connection will be accepted. Otherwise if the count = 0, the server will reject this connection.
 - 3, When a connection is closed, the server increment the semaphore, which means a slot has been freed and the server now can accept a new connection.
-

Problem 9 (20 points) Write a multithreaded program that outputs prime numbers. This program should work as follows: The user will run the program and will enter a number on the command line. The program will then create a separate thread that outputs all the prime numbers less than or equal to the number entered by the user.

Answer: the source code:

```
#include <iostream>
#include <thread>
#include <vector>
#include <cmath>

// Function to check if a number is prime
bool is_prime(int num) {
    if (num <= 1) return false;
    if (num <= 3) return true;
    if (num % 2 == 0 || num % 3 == 0) return false;
    for (int i = 5; i * i <= num; i += 6) {
        if (num % i == 0 || num % (i + 2) == 0) return false;
    }
    return true;
}

// Worker function that prints prime numbers
void print_primes(int limit) {
    for (int num = 2; num <= limit; ++num) {
        if (is_prime(num)) {
            std::cout << num << " ";
        }
    }
}
```

```

    }
}
std::cout << std::endl;
}

int main(int argc, char* argv[]) {
    if (argc != 2) {
        std::cerr << "Usage: " << argv[0] << " <number>" << std::endl;
        return 1;
    }

    int limit;
    try {
        limit = std::stoi(argv[1]);
    } catch (const std::invalid_argument& ia) {
        std::cerr << "Please enter a valid integer." << std::endl;
        return 1;
    } catch (const std::out_of_range& oor) {
        std::cerr << "Number out of range." << std::endl;
        return 1;
    }

    // Create a thread to print prime numbers
    std::thread prime_thread(print_primes, limit);
    prime_thread.join(); // Wait for the thread to finish

    return 0;
}

```

The printed result:

```

● kevin@Kevin:/mnt/c/Github/coding-interview-university/OperatingSystem_Linux/CSE-5305$ g++ -o print_prime print_prime.cpp
● kevin@Kevin:/mnt/c/Github/coding-interview-university/OperatingSystem_Linux/CSE-5305$ ./print_prime 100
2 3 5 7 11 13 17 19 23 29 31 37 41 43 47 53 59 61 67 71 73 79 83 89 97

```